

| | |
|---|---|
| <pre> ML.Uber1 import pandas as pd df = pd.read_csv('uber.csv') df.head();; <pre> <i>*pre-process the dataset</i> df.info();; df.describe();; df.isnull().sum();; print(df.isnull().values.sum());; df = df.drop(['Unnamed: 0', 'key'], axis = 1);; df.head();; df = df[df.fare_amount > 0];; df.shape();; df.describe();; df = df[(df.passenger_count <= 6) & (df.passenger_count > 0)];; df.head();; df = df[(df.pickup_longitude.between(-180,180,inclusive = "both")) & (df.pickup_latitude.between(-90,90,inclusive = "both")) & (df.dropoff_longitude.between(- 180,180,inclusive = "both")) & (df.dropoff_latitude.between(-90,90,inclusive = "both"))];; df.head(10);; df.info();; df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime']);; df.info();; df.head(10);; import calendar;; df['year'] = df.pickup_datetime.dt.year df['month'] = df.pickup_datetime.dt.month df['weekday'] = df.pickup_datetime.dt.weekday df['hour'] = df.pickup_datetime.dt.hour;; df = df.drop(['pickup_datetime'], axis=1);; df.head(10);; df.describe();; <i>*Remove outliers</i> df = df.reset_index();; df.head();; df = df.drop(['index'], axis = 1);; df.head(10);; import numpy as np;; Q1 = np.percentile(df['fare_amount'], 25) Q3 = np.percentile(df['fare_amount'], 75) IQR = Q3 - Q1 print(f'IQR = {IQR}');; upper = np.where(df['fare_amount'] >= (Q3 + 1.5*IQR)) lower = np.where(df['fare_amount'] <= (Q1 - 1.5*IQR)) print(f'upper = {upper}') print(f'lower = {lower}');; df = df.drop(upper[0]);; df = df.drop(lower[0]);; df.describe();; <i>*Find Correlation</i> import seaborn as sns import matplotlib.pyplot as plt uber_corr = df.corr() #use heatmap plt.figure(figsize=(10,7)) sns.heatmap(uber_corr,annot=True) plt.show();; <i>*implement linear reg. & random fore reg.</i> from sklearn.model_selection import train_test_split;; X = df.drop('fare_amount', axis = 1) y = df["fare_amount"];; X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2);; from sklearn.linear_model import LinearRegression lrmodel = LinearRegression() lrmodel.fit(X_train, y_train) lr_pred = lrmodel.predict(X_test) lr_pred;; from sklearn.ensemble import RandomForestRegressor rfmodel = RandomForestRegressor(n_estimators=100, random_state=101) rfmodel.fit(X_train, y_train) rf_pred = rfmodel.predict(X_test) rf_pred;; <i>*evaluate model</i> from sklearn.metrics import mean_squared_error;; lrmodel_rmse = np.sqrt(mean_squared_error(y_test, lr_pred)) rfmodel_rmse = np.sqrt(mean_squared_error(y_test, rf_pred)) print(f'Linear Regression RMSE = {lrmodel_rmse}') print(f'Random Forest RMSE = {rfmodel_rmse}');; from sklearn.metrics import r2_score lrmodel_r2 = r2_score(y_test, lr_pred) rfmodel_r2 = r2_score(y_test, rf_pred) print(f'Linear Regression R2 = {lrmodel_r2}') </pre> </pre> | <pre> ML3GradientDescentAlgorithm import pandas as pd cur_x = 2 rate = 0.01 precision = 0.0000001 previous_step_size = 1 max_iters = 10000 iters = 0 df = lambda x: (2 * (x + 3));; while previous_step_size > precision and iters < max_iters: prev_x = cur_x cur_x = cur_x - rate * df(prev_x) previous_step_size = abs(cur_x - prev_x) iters = iters + 1 print(f'Iteration {iters} \n value is {cur_x}') print(f'The local minima occurs at {cur_x}') </pre> <pre> ML2bankcustomer,build a neuralnetwork import pandas as pd import numpy as np;; ds = pd.read_csv('Churn_Modelling.csv') ds.head(10);; ds.columns;; ds.shape;; ds['Geography'].value_counts(normalize=True);; ds = ds.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1);; ds.info();; ds.describe();; X = ds.iloc[:,0:10].values y = ds.iloc[:,10].values;; X;; from sklearn.preprocessing import LabelEncoder print(X[:8,1], '... will now become: ') label_X_country_encoder = LabelEncoder() X[:,1] = label_X_country_encoder.fit_transform(X[:,1]) print(X[:8,1]);; print(X[:6,2], '... will now become: ') label_X_gender_encoder = LabelEncoder() X[:,2] = label_X_gender_encoder.fit_transform(X[:,2]) print(X[:6,2]);; X.shape;; from sklearn.compose import ColumnTransformer from sklearn.preprocessing import OneHotEncoder transform = ColumnTransformer([("countries", OneHotEncoder(), [1])], remainder="passthrough") X = transform.fit_transform(X);; X;; X = X[:,1:];; X.shape;; from sklearn.model_selection import train_test_split X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.2, random_state = 0) numeric_cols = ['CreditScore','Age','Tenure','Balance','NumOfProducts','EstimatedSalary'];; X_train[:,np.array([2,4,5,6,7,10])];; from sklearn.preprocessing import StandardScaler sc=StandardScaler() X_train[:,np.array([2,4,5,6,7,10])] = sc.fit_transform(X_train[:,np.array([2,4,5,6,7,10])]) X_test[:,np.array([2,4,5,6,7,10])] = sc.transform(X_test[:,np.array([2,4,5,6,7,10])]);; X_train[0];; from sklearn.preprocessing import StandardScaler sc=StandardScaler() X_train = sc.fit_transform(X_train) X_test = sc.transform(X_test);; X_train;; X_train.shape;; from tensorflow.keras.models import Sequential # Initializing the ANN classifier = Sequential();; from tensorflow.keras.layers import Dense classifier.add(Dense(activation = 'relu', input_dim = 11, units=16, kernel_initializer='uniform'));; classifier.add(Dense(8, activation='relu', kernel_initializer='uniform'));; classifier.add(Dense(1, activation = 'sigmoid', kernel_initializer='uniform'));; 192-11*16;; classifier.summary();; </pre> |
|---|---|

```
print(f'Random Forest R2 = {rfmodel_r2}');
```

ML5KMeansClustering

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt;
from sklearn.cluster import KMeans, k_means
from sklearn.decomposition import PCA;
df = pd.read_csv('/content/sales_data_sample.csv');
df.head();
df.shape;
df.describe();
df.info();
df.isnull().sum();
df.dtypes;
df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE', 'CITY', 'TERRITORY',
'PHONE', 'STATE', 'CONTACTFIRSTNAME', 'CONTACTLASTNAME', 'CUSTOMERNAME',
'ORDERNUMBER']
df = df.drop(df_drop, axis=1);
df.isnull().sum();
df.dtypes;
*checking the categorical columns
df['COUNTRY'].unique();
df['PRODUCTLINE'].unique();
df['DEALSIZE'].unique();
productline = pd.get_dummies(df['PRODUCTLINE'])
Dealsize = pd.get_dummies(df['DEALSIZE']);
df = pd.concat([df,productline,Dealsize], axis = 1);
df_drop = ['COUNTRY', 'PRODUCTLINE', 'DEALSIZE'] #Dropping Country too as there are alot of
countries.
df = df.drop(df_drop, axis=1);
df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes;
df.drop('ORDERDATE', axis=1, inplace=True);
df.dtypes;
distortions = []
K = range(1,10)
for k in K:
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(df)
    distortions.append(kmeanModel.inertia_);
plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k')
plt.show();
*numb. Of k increases inertia decreases
X_train = df.values;
X_train.shape;
model = KMeans(n_clusters=3, random_state=2)
model = model.fit(X_train)
predictions = model.predict(X_train);
unique_counts = np.unique(predictions, return_counts=True);
counts = counts.reshape(1,3);
counts_df = pd.DataFrame(counts, columns=['Cluster1', 'Cluster2', 'Cluster3']);
counts_df.head();
pca = PCA(n_components=2);
reduced_X = pd.DataFrame(pca.fit_transform(X_train), columns=['PCA1', 'PCA2']);
reduced_X.head();
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'], reduced_X['PCA2']);
model.cluster_centers_;
reduced_centers = pca.transform(model.cluster_centers_);
reduced_centers;
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'], reduced_X['PCA2'])
plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=300);
reduced_X['Clusters'] = predictions;
reduced_X.head();
plt.figure(figsize=(14,10))

plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters']
== 0].loc[:, 'PCA2'], color='slateblue') .
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters']
== 1].loc[:, 'PCA2'], color='springgreen') .
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters']
== 2].loc[:, 'PCA2'], color='indigo') .

plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=300);
```

```
classifier.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy']);
classifier.summary();
classifier.fit(X_train, y_train,
validation_data=(X_test, y_test),
epochs=20);
y_pred = classifier.predict(X_test)
print(y_pred);
y_pred = (y_pred > 0.5)
print(y_pred);
from sklearn.metrics import confusion_matrix, classification_report

cm1 = confusion_matrix(y_test, y_pred)
print(cm1);
print(classification_report(y_test, y_pred));
accuracy_model1 =
((cm1[0][0]+cm1[1][1])*100)/(cm1[0][0]+cm1[0][1]+cm1[1][1]+cm1[1][0])
print (accuracy_model1, '% of testing data was classified correctly');
classifier.summary();
classifier.compile(optimizer='adam', loss = 'binary_crossentropy',
metrics=['accuracy']);
classifier.fit(X_train, y_train,
validation_data=(X_test, y_test),
epochs=20,
batch_size=32);
y_pred = classifier.predict(X_test)
print(y_pred);
y_pred = (y_pred > 0.5)
print(y_pred);
cm2 = confusion_matrix(y_test, y_pred)
print(cm2);
cm2 = classification_report(y_test, y_pred)
print(cm2);
```

ML4KNN

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import metrics;
df=pd.read_csv('/content/diabetes.csv');
df.columns;
df.isnull().sum();
X = df.drop('Outcome', axis = 1)
y = df['Outcome'];
from sklearn.preprocessing import scale
X = scale(X)
# split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state =
42);
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)

knn.fit(X_train, y_train)
y_pred = knn.predict(X_test);
print("Confusion matrix: ")
cs = metrics.confusion_matrix(y_test, y_pred)
print(cs);
print("Accuracy ", metrics.accuracy_score(y_test, y_pred));
total_misclassified = cs[0,1] + cs[1,0]
print(total_misclassified)
total_examples = cs[0,0]+cs[0,1]+cs[1,0]+cs[1,1]
print(total_examples)
print("Error rate", total_misclassified/total_examples)
print("Error rate ", 1-metrics.accuracy_score(y_test, y_pred));
print("Precision score", metrics.precision_score(y_test, y_pred));
print("Recall score ", metrics.recall_score(y_test, y_pred));
print("Classification report ", metrics.classification_report(y_test, y_pred));
```

